

UC BERKELEY CS 185/285 • PRIORITY 1 COMPLETE MASTER COMPENDIUM

Robotics Reinforcement Learning Foundations

A Rigorous, Pedagogical Manual for Dual-Arm Robotic Manipulation & Contact Control (Paper 1: Adaptive Soft Fruit Slicing with PPO / SAC in NVIDIA Isaac Lab via SkRL)

Lectures Unified: CS285 L01, L04, L05, L06, L08, L10
Instructor: Prof. Sergey Levine (UC Berkeley RAIL Lab)

Research Lab: DEX-ROB Lab, Tianjin University (Prof. Shan An)

Environment & Stack: NVIDIA Isaac Lab / PhysX 5 • SkRL Framework • PyTorch

🔗 HOW THIS MASTER GUIDE ORGANIZES PRIORITY 1 CURRICULUM

This master reference document synthesizes the six core lectures of UC Berkeley CS285 into an integrated, end-to-end framework. It bridges theoretical foundations (Lectures 1 & 4), direct policy optimization (Lecture 5), actor-critic variance reduction (Lecture 6), off-policy continuous control (Lecture 8), and trust-region stabilization (Lecture 10). It culminates in a complete operational blueprint for your dual-arm robotic slicing task, the 12-dimension algorithm comparison matrix, and a 10-question defense cheatsheet.

Module 1: The Foundations & The Robotics Closed Loop (Lecture 1)

(CS285 Lecture 1 • Video ID: DD8APgTEix4) Welcome to the mathematical study of sequential decision making. In traditional robotics, control laws are derived from explicit kinematic and dynamic models ($F = Ma$). However, when interacting with non-linear, viscoelastic, fracturing materials (such as ripe fruits), analytical models fail.

1.1 Why Reinforcement Learning Over Supervised Imitation?

In **Supervised Learning** (Behavioral Cloning), an expert teleoperates the robot and records state-action pairs (s, a) . The network minimizes $\mathcal{L}_{BC}(\theta) = \mathbb{E}_{(s,a) \sim \mathcal{D}_{\text{expert}}} [\|a - \pi_{\theta}(s)\|^2]$. While intuitive, this approach suffers from a catastrophic mathematical vulnerability: **Covariate Shift**.

📌 ROSS & BAGNELL (2011) COMPOUNDING ERROR PROOF

Let the probability of the learned policy making an ϵ -mistake on the expert training distribution be ϵ :

$$\mathbb{P}_{s \sim d_{\pi^*}}(\pi_{\theta}(s) \neq \pi^*(s)) \leq \epsilon$$

Under open-loop behavioral cloning, if an error occurs at time step t , the robot enters an unvisited state space $s \notin \mathcal{D}_{\text{expert}}$. Because the policy was never trained on recovery actions, all subsequent actions can fail. The expected total trajectory error over time horizon T compounds quadratically:

$$\mathbb{E}[\text{Errors}_{1:T}] \leq \sum_{t=1}^T \epsilon \cdot (T - t) \approx \frac{1}{2} \epsilon T^2 = \mathcal{O}(\epsilon T^2)$$

In contrast, **Reinforcement Learning operates in closed-loop**, exploring its own induced distribution $d^{\pi_{\theta}}(s)$. When the knife slips or the fruit deforms unpredictably, the agent actively learns corrective policies, bounding error strictly to $\mathcal{O}(\epsilon T)$.



1.2 The Credit Assignment Problem & Dense Shaping

Prof. Levine highlights the fundamental challenge of **Credit Assignment**: when an episode consists of 500 robot control steps, and the tomato bursts at step 480, which specific millisecond torque spike was responsible?

💡 PLAIN ENGLISH: THE EXAM ANALOGY

If you study all semester and receive an **F** on the final exam, that single terminal letter grade does not tell you whether you failed because you slept late on Tuesday or because you misunderstood Chapter 3. In RL, sparse terminal rewards require millions of sample trajectories to backpropagate credit. We solve this through **Dense Progress Shaping** and **Value Functions** that score state quality at every microsecond.

Module 2: The Core Mathematics & MDP Formulation (Lecture 4)

(CS285 Lecture 4 • Video ID: Fcplul7rAEE) Lecture 4 establishes the formal mathematical foundation of Reinforcement Learning. Every physical quantity in the robotics laboratory corresponds to a precise mathematical object.

2.1 The Markov Decision Process: $\mathcal{M} = \langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma, \rho_0 \rangle$

Symbol	Formal Math	Your Paper 1 Physical Translation
State (s)	$s_t \in \mathcal{S}$	33 numbers: Dual-arm joint positions & velocities (14), knife pose & velocities (9), tomato pose (7), TacBlade F/T (6), acoustic burst (1).
Action (a)	$a_t \in \mathcal{A}$	6 numbers: Downward feed delta Δv_z , lateral sawing speed v_{slice} , vertical stiffness delta ΔK , damping delta ΔD , holding arm stabilization.
Dynamics ($\mathcal{P}$)	$P(s_{t+1} \mid s_t, a_t)$	Isaac Sim PhysX 5 Engine: Computes knife penetration, tomato deformation, friction, and outputs the next sensor state.
Reward ($\mathcal{R}$)	$r(s_t, a_t)$	Scoring function: Penetration progress + Sawing bonus - Crushing force penalty (>8N) - Jitter penalty.
Discount (γ)	$\gamma = 0.99$	Patience factor: Ensures the policy plans for the complete cut rather than stopping on first skin contact.

2.2 The Markov Property & Trajectory Probability Factorization

A stochastic process satisfies the **Markov Property** if the conditional probability distribution of future states depends only upon the current state and action, and is conditionally independent of all historical states:

$$\mathbb{P}(s_{t+1} \mid s_t, a_t, s_{t-1}, a_{t-1}, \dots, s_0, a_0) = \mathbb{P}(s_{t+1} \mid s_t, a_t)$$

By applying the probability chain rule, the probability distribution of an entire trajectory $\tau = (s_0, a_0, s_1, a_1, \dots, s_T)$ factorizes cleanly:

$$\mathbb{P}(\tau \mid \theta) = \rho_0(s_0) \prod_{t=0}^{T-1} \pi_\theta(a_t \mid s_t) P(s_{t+1} \mid s_t, a_t)$$

2.3 The Four Core Value Functions

RL defines four fundamental value functions that measure future performance:

1. **State-Value Function** $V^\pi(s)$: Expected cumulative return starting from state s following policy π :

$$V^\pi(s) = \mathbb{E}_\pi \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) \mid s_0 = s \right]$$

2. **Action-Value Function** $Q^\pi(s, a)$: Expected return starting from s , taking arbitrary action a , and thereafter following π :

$$Q^\pi(s, a) = r(s, a) + \gamma \mathbb{E}_{s' \sim P(\cdot \mid s, a)} [V^\pi(s')]$$

3. **Optimal State-Value Function** $V^*(s)$: Maximum attainable return across all conceivable policies:

$$V^*(s) = \max_{\pi} V^{\pi}(s) = \max_{a \in \mathcal{A}} Q^*(s, a)$$

4. **Advantage Function** $A^{\pi}(s, a) = Q^{\pi}(s, a) - V^{\pi}(s)$: The relative superiority of taking action a compared to the policy's average expected outcome in state s .
- If $A > 0$: Action outperformed expectations $\rightarrow$ Increase selection probability.
 - If $A < 0$: Action underperformed $\rightarrow$ Decrease selection probability.

BELLMAN CONTRACTION OPERATOR & BANACH FIXED POINT PROOF

Define the Bellman optimality backup operator $\mathcal{B} : \mathbb{R}^{|\mathcal{S}|} \rightarrow \mathbb{R}^{|\mathcal{S}|}$ on value vector V :

$$(\mathcal{B}V)(s) = \max_{a \in \mathcal{A}} \left[r(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s' | s, a) V(s') \right]$$

For any two arbitrary value functions u and v , evaluate the infinity norm difference $\|\mathcal{B}u - \mathcal{B}v\|_{\infty}$:

$$|(\mathcal{B}u)(s) - (\mathcal{B}v)(s)| \leq \gamma \max_{a \in \mathcal{A}} \sum_{s'} P(s' | s, a) |u(s') - v(s')| \leq \gamma \|u - v\|_{\infty} \sum_{s'} P(s' | s, a) = \gamma \|u - v\|_{\infty}$$

Taking the supremum over all states $s \in \mathcal{S}$ proves that $\|\mathcal{B}u - \mathcal{B}v\|_{\infty} \leq \gamma \|u - v\|_{\infty}$. Since $\gamma \in [0, 1)$, $\mathcal{B}$ is a **strict γ -contraction mapping**. By the *Banach Fixed Point Theorem*, repeated application $\lim_{k \rightarrow \infty} \mathcal{B}^k V_0$ converges uniquely to the optimal value function V^* , independent of initial guess V_0 .

Module 3: Direct Policy Optimization & Policy Gradients (Lecture 5)

(CS285 Lecture 5 • Video ID: S0D9REIVdg4) How do we optimize neural network policy weights θ when the physical environment contains non-differentiable contact fractures? This is the core domain of **Policy Gradient Methods**.

3.1 The Fundamental Dilemma: Non-Smooth Contact Dynamics

! WHY ANALYTICAL BACKPROPAGATION THROUGH PHYSICS FAILS

A tomato skin tearing under a sharp blade is an irreversible physical rupture. Dynamic friction switches discontinuously between stick and slip regimes. Attempting to backpropagate gradients $\frac{\partial s_{t+1}}{\partial a_t}$ through rigid-body contact solvers results in either zero gradients (flat plateaus) or infinite gradients (delta spikes). Policy gradients bypass this by treating the environment dynamics as a black box.

3.2 Derivation of the Policy Gradient Theorem

Let the expected total discounted return under policy π_θ be $J(\theta) = \mathbb{E}_{\tau \sim \pi_\theta}[R(\tau)] = \int \mathbb{P}(\tau | \theta) R(\tau) d\tau$. We differentiate directly:

$$\nabla_\theta J(\theta) = \int \nabla_\theta \mathbb{P}(\tau | \theta) R(\tau) d\tau$$

Applying the **Log-Derivative Trick** ($\nabla_\theta f(\theta) = f(\theta) \nabla_\theta \log f(\theta)$):

$$\nabla_\theta J(\theta) = \int \mathbb{P}(\tau | \theta) \nabla_\theta \log \mathbb{P}(\tau | \theta) R(\tau) d\tau = \mathbb{E}_{\tau \sim \pi_\theta} [\nabla_\theta \log \mathbb{P}(\tau | \theta) R(\tau)]$$

Substituting the trajectory log-probability $\log \mathbb{P}(\tau | \theta) = \log \rho_0(s_0) + \sum_{t=0}^{T-1} \log \pi_\theta(a_t | s_t) + \sum_{t=0}^{T-1} \log P(s_{t+1} | s_t, a_t)$, and noting that the initial distribution and transition dynamics do not depend on θ :

$$\nabla_\theta \log \mathbb{P}(\tau | \theta) = \sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t | s_t)$$

$$\nabla_\theta J(\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[\sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t | s_t) \cdot \sum_{t'=0}^{T-1} r(s_{t'}, a_{t'}) \right]$$

3.3 Continuous Gaussian Policy Score Function

In dual-arm continuous impedance control, actions are continuous real vectors $a \sim \mathcal{N}(\mu_\theta(s), \Sigma_\theta(s))$. For a diagonal covariance matrix with σ_i :

$$\log \pi_\theta(a | s) = -\frac{d}{2} \log(2\pi) - \sum_{i=1}^d \log \sigma_i - \frac{1}{2} \sum_{i=1}^d \frac{(a_i - \mu_i(s))^2}{\sigma_i^2}$$

$$\nabla_\mu \log \pi_\theta(a | s) = \frac{a - \mu_\theta(s)}{\sigma^2}$$

💡 PHYSICAL INTERPRETATION OF GAUSSIAN POLICY GRADIENTS

Notice the elegance of the gradient: $\frac{a - \mu_\theta(s)}{\sigma^2} \cdot \hat{Q}$. If a trial action a was higher than the mean μ , and achieved positive advantage ($\hat{Q} > 0$), the product is positive, pulling the mean μ closer to a . If the action resulted in crushed pulp ($\hat{Q} < 0$), the product is negative, pushing the mean in the opposite direction.

3.4 Essential Variance Reduction: Causality & Optimal Baselines

Raw REINFORCE suffers from excessive variance. Modern implementations enforce two mandatory structural transformations:

1. **Causality ("Reward-to-Go"):** Actions taken at step t cannot influence past rewards $r_{t'}$ for $t' < t$. Past terms have expectation zero and are safely dropped:

$$\nabla_\theta J(\theta) = \mathbb{E} \left[\sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t | s_t) \sum_{t'=t}^{T-1} \gamma^{t'-t} r(s_{t'}, a_{t'}) \right]$$

2. **Expected Grad-Log-Prob (EGLP) Lemma & Baselines:** Subtracting any state-dependent baseline $b(s)$ does not introduce bias:

$$\mathbb{E}_{a \sim \pi} [\nabla_\theta \log \pi_\theta(a | s) b(s)] = b(s) \int \pi_\theta(a | s) \frac{\nabla_\theta \pi_\theta(a | s)}{\pi_\theta(a | s)} da = b(s) \nabla_\theta \int \pi_\theta(a | s) da = b(s) \nabla_\theta (1) = 0$$

Module 4: Actor-Critic Architectures & GAE (Lecture 6)

(CS285 Lecture 6 • Video ID: MzIWnZrCvw) While baselines reduce variance, Monte Carlo returns still accumulate noise across long horizons. **Actor-Critic** architectures replace full rollouts with a learned Critic network.

4.1 The Dual-Network Topology



4.2 Temporal Difference (TD) Bootstrapping & The Semi-Gradient Detach Rule

The Critic approximates $V(s)$ by minimizing the mean squared Bellman error against a bootstrapped 1-step target:

$$\mathcal{L}(\phi) = \frac{1}{2} \mathbb{E}_{(s_t, a_t, r_t, s_{t+1})} \left[(r_t + \gamma V_{\phi}(s_{t+1}) - V_{\phi}(s_t))^2 \right]$$

● CRUCIAL PYTORCH RULE: SEMI-GRADIENT DETACH

When backpropagating through the Critic loss, you must **detach** the target: `target = reward + gamma * critic(next_state).detach()`. If you do not call `.detach()`, the optimizer updates the target network concurrently with the current value estimate, creating an unstable feedback loop where the network chases its own tail.

4.3 Generalized Advantage Estimation (GAE- λ) Complete Derivation

Define the 1-step TD residual error: $\delta_t^V = r_t + \gamma V(s_{t+1}) - V(s_t)$. Now observe the k -step advantage estimators:

$$\begin{aligned} \hat{A}_t^{(1)} &= \delta_t^V = r_t + \gamma V(s_{t+1}) - V(s_t) \\ \hat{A}_t^{(2)} &= \delta_t^V + \gamma \delta_{t+1}^V = r_t + \gamma r_{t+1} + \gamma^2 V(s_{t+2}) - V(s_t) \\ \hat{A}_t^{(k)} &= \sum_{l=0}^{k-1} \gamma^l \delta_{t+l}^V = \sum_{l=0}^{k-1} \gamma^l r_{t+l} + \gamma^k V(s_{t+k}) - V(s_t) \end{aligned}$$

GAE takes an exponentially decaying weighted average of all k -step estimators using parameter $\lambda \in [0, 1]$:

$$\hat{A}_t^{\text{GAE}(\gamma, \lambda)} = (1 - \lambda) \sum_{k=1}^{\infty} \lambda^{k-1} \hat{A}_t^{(k)} = \sum_{l=0}^{\infty} (\gamma \lambda)^l \delta_{t+l}^V$$

⚙️ SKRL TUNING: THE $\lambda = 0.95$ SWEET SPOT

Setting $\lambda = 0$ yields pure TD(0) (minimal variance, but high bias from an imperfect critic). Setting $\lambda = 1$ yields pure Monte Carlo (unbiased, but noisy). In your SkRL configuration, `gae_lambda: 0.95` captures 95% of the variance-reduction benefits of the Critic while eliminating persistent bootstrapping bias.

Module 5: Continuous Control & Soft Actor-Critic (Lecture 8)

(CS285 Lecture 8 • Video ID: IQaVa53pS-Q) In discrete environments like Atari, Q-learning calculates $\arg \max_a Q(s, a)$ by simple enumeration. On continuous robotic arms, finding this maximum is a complex non-linear optimization problem.

5.1 The Continuous Action Maximization Crisis

On a robot arm with 6-DoF continuous velocity commands $a \in [-1, 1]^6$, evaluating $\max_a Q(s, a)$ across infinite real vectors at 500 Hz is intractable. Discretization scales exponentially (K^6), while sampling methods (Cross-Entropy Method) are too slow for real-time control.

The Modern Solution (Actor Maximizer): Train an explicit Actor network $\pi_\theta(s)$ parameterized to directly output the action that maximizes the Critic: $\max_\theta \mathbb{E}_{s \sim \mathcal{D}}[Q_\phi(s, \pi_\theta(s))]$.

5.2 Sutton's Deadly Triad & Stabilization Mechanisms

Combining **Function Approximation** (Deep Nets), **Bootstrapping** (Bellman TD targets), and **Off-Policy Training** (Replay Buffer) leads to the *Deadly Triad*, often causing value estimates to diverge to infinity. Modern algorithms tame this via three structural pillars:

Mechanism	Mathematical Formulation	Physical Purpose in Robotics
Polyak Soft Updates	$\bar{\theta} \leftarrow \tau \theta + (1 - \tau) \bar{\theta}$ with $\tau = 0.005$	Creates a stable, slow-moving target network with half-life $t_{1/2} = \frac{\ln 2}{\tau} \approx 138$ steps, preventing gradient explosion.
Clipped Twin-Q (Double Q)	$y = r + \gamma \min(Q_{\bar{\theta}_1}(s', a'), Q_{\bar{\theta}_2}(s', a'))$	Overcomes positive maximization bias caused by Jensen's inequality: $\mathbb{E}[\max(X_1, X_2)] \geq \max(\mathbb{E}[X_1], \mathbb{E}[X_2])$.
Experience Replay	Sample $(s, a, r, s') \sim \mathcal{D}$	Breaks high temporal correlation across sequential simulator frames, stabilizing gradient descent.

5.3 Soft Actor-Critic (SAC) & Maximum Entropy Objective

Standard RL maximizes expected reward. **Maximum Entropy RL** augments the reward with the Shannon entropy $\mathcal{H}(\pi(\cdot | s_t)) = \mathbb{E}_{a \sim \pi}[-\log \pi(a | s_t)]$ of the policy:

$$J(\pi) = \sum_{t=0}^T \mathbb{E}_{(s_t, a_t) \sim \rho_\pi} [r(s_t, a_t) + \alpha \mathcal{H}(\pi(\cdot | s_t))]$$

💡 PLAIN ENGLISH: PAYING THE ROBOT TO STAY CURIOUS

In soft tomato slicing, an agent maximizing only reward easily gets trapped in a local minimum: it hovers the knife 1 mm above the tomato skin to avoid crushing penalties. The entropy temperature α rewards the robot for testing multiple blade orientations and speeds, discovering successful puncture techniques.

5.4 The Reparameterization Trick & Tanh Squashing Jacobian Correction

To differentiate through stochastic actions, SAC samples Gaussian noise $\epsilon \sim \mathcal{N}(0, I)$ and transforms it: $u = \mu_\theta(s) + \sigma_\theta(s) \odot \epsilon$, bounded to $[-1, 1]$ via $a = \tanh(u)$. Because $\tanh$ is non-linear, the probability density requires a Jacobian determinant correction:

$$\log \pi(a \mid s) = \log \mu(u \mid s) - \sum_{i=1}^d \log(1 - \tanh^2(u_i))$$

Module 6: Advanced Policy Gradients & PPO (Lecture 10)

(CS285 Lecture 10 • Video ID: m7IU5KBS4sw) Proximal Policy Optimization (PPO) is the de-facto standard for GPU-accelerated robotics. This module explores why standard policy gradients suffer from policy collapse and how PPO guarantees monotonic policy improvement.

6.1 The Catastrophic Failure Mode: Policy Collapse

☀ WHY POLICY COLLAPSE OCCURS IN REINFORCEMENT LEARNING

In supervised learning, an excessively large learning rate causes a momentary loss spike, but training recovers on subsequent batches. In RL, **the policy generates its own training distribution**. If an overly aggressive gradient step drives policy parameters into an unstable region, the robot executes erratic motions. The resulting trajectory batch contains zero successful cuts. Subsequent updates compute gradients on degenerate data, initiating an unrecoverable death spiral.

6.2 Kakade-Langford Monotonic Improvement Theorem

Kakade & Langford (2002) proved that the performance difference between two arbitrary policies $\tilde{\pi}$ and π is exactly governed by:

$$\eta(\tilde{\pi}) - \eta(\pi) = \mathbb{E}_{\tau \sim \tilde{\pi}} \left[\sum_{t=0}^{\infty} \gamma^t A^{\pi}(s_t, a_t) \right]$$

Because sampling from the new policy $\tilde{\pi}$ during optimization is impossible, we optimize the surrogate objective $L_{\pi}(\tilde{\pi})$ using importance sampling, subject to a theoretical lower bound:


$$\eta(\tilde{\pi}) \geq L_{\pi}(\tilde{\pi}) - C \cdot \max_s D_{\text{KL}}(\pi(\cdot | s) \| \tilde{\pi}(\cdot | s)) \quad \text{where} \quad C = \frac{2\epsilon\gamma}{(1-\gamma)^2}$$

6.3 The PPO Clipped Surrogate Objective

While TRPO enforces the KL constraint using second-order natural gradients (Fisher Information Matrix), PPO enforces the trust region using a first-order clipped surrogate loss. Define the importance sampling probability ratio:

$$r_t(\theta) = \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_{\text{old}}}(a_t | s_t)}$$

$$L^{\text{CLIP}}(\theta) = \hat{\mathbb{E}}_t \left[\min \left(r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t \right) \right]$$



6.4 Complete 4-Quadrant PPO Dynamics

Advantage	Ratio Condition	Objective Value	Optimization Effect
$\hat{A} > 0$ (Good)	$r_t(\theta) \leq 1 + \epsilon$	$r_t(\theta)\hat{A}_t$	Standard Positive Gradient: Increases probability of successful knife motion.
$\hat{A} > 0$ (Good)	$r_t(\theta) > 1 + \epsilon$	$(1 + \epsilon)\hat{A}_t$	Clipped (Zero Gradient): Prevents policy from over-committing to an action.
$\hat{A} < 0$ (Bad)	$r_t(\theta) \geq 1 - \epsilon$	$r_t(\theta)\hat{A}_t$	Standard Negative Gradient: Suppresses crushing action.
$\hat{A} < 0$ (Bad)	$r_t(\theta) < 1 - \epsilon$	$(1 - \epsilon)\hat{A}_t$	Clipped (Zero Gradient): Prevents destructive gradient updates when probability is already low.

Module 7: The Master 12-Dimension Algorithm Grand Comparison Matrix

This authoritative reference matrix compares the six seminal reinforcement learning algorithms across twelve fundamental architectural, mathematical, and practical dimensions.

Dimension	REINFORCE	A2C	DQN	SAC	TRPO	PPO
1. Paradigm	On-Policy PG	On-Policy AC	Off-Policy TD	Off-Policy MaxEnt AC	On-Policy Trust Region	On-Policy Clipped AC
2. Policy Type	Stochastic $\pi_{\theta}(a s)$	Stochastic $\pi_{\theta}(a s)$	Implicit $\arg \max Q$	Stochastic Gaussian	Stochastic Gaussian	Stochastic Gaussian
3. Action Space	Discrete / Cont.	Discrete / Cont.	Discrete only	Continuous only	Discrete / Cont.	Discrete / Cont.
4. Sample Efficiency	Very Low (10^7)	Low (10^7)	High (10^5)	Very High (10^5)	Moderate (10^6)	Moderate (10^6)
5. Wall-Clock Speed	Slow	Fast	Moderate	Moderate (Replay)	Slow (Fisher Hessian)	Fastest on GPU
6. Objective Function	$\mathbb{E}[\nabla \log \pi G_t]$	$\mathbb{E}[\nabla \log \pi \hat{A}^{\text{TD}}]$	$\mathbb{E}[(r + \gamma \max Q - Q)^2]$	$\mathbb{E}[r + \alpha \mathcal{H}(\pi)]$	$\max L_{\pi} \text{ s.t. } D_{\text{KL}} \leq \delta$	$\mathbb{E}[\min(r\hat{A}, \text{clip})]$
7. Stability	Extremely Low	Moderate	Moderate	High (Twin-Q + Polyak)	Very High	Very High
8. Replay Buffer	No	No	Yes	Yes	No	No
9. Exploration	Stochastic policy	Entropy bonus	ϵ -greedy	MaxEnt ($\alpha \mathcal{H}$)	Stochastic policy	Entropy loss scaling
10. Sim-to-Real	Poor	Moderate	Poor (Continuous)	High	High	Highest (Dominant)
11. Parallel Sim	Poor	Good	Poor (Async buffer)	Moderate	Poor (CG overhead)	Optimal (Linear)
12. Canonical Use	Toy CartPole	Atari 2600	Video Games	MuJoCo Continuous	Robotics Locomotion	Isaac Lab / Manipulation

Module 8: Master Paper 1 Implementation Blueprint (Dual-Arm Slicing)

Here is the comprehensive architectural and operational blueprint connecting CS285 Priority 1 directly to your Master's thesis methodology: **Adaptive Dual-Arm Soft Object Slicing with PPO / SAC in NVIDIA Isaac Lab via SkRL**.

8.1 System Specifications & Multi-Modal Sensory Integration

The robotic slicing setup consists of two collaborative arms: a **Holding Arm** that stabilizes the soft tomato with gentle grasp force, and a **Cutting Arm** equipped with a sensorized blade containing an **ATI Nano17 6-axis F/T transducer** and a **piezoelectric acoustic burst sensor**.

Subsystem	Dimension	Sensor / Physical Origin	Normalization & Noise Model
Arm Kinematics	$\mathbb{R}^{14}$	Dual-arm joint positions q (7) and velocities $\dot{q}$ (7).	Normalized by joint limits; Gaussian noise $\sigma = 0.005$ rad.
Blade Kinematics	$\mathbb{R}^9$	Blade tip 3D position p , linear velocity v , angular velocity ω .	Relative to cutting board frame; $\sigma_v = 0.01$ m/s.
Tomato Geometry	$\mathbb{R}^7$	3D center of mass p_{fruit} and orientation quaternion q_{fruit} .	Sim: USD pose; Real: YOLOv8 + Depth camera.
TacBlade F/T	$\mathbb{R}^6$	3-axis contact force (F_x, F_y, F_z) and torque (T_x, T_y, T_z) .	Filtered with 50 Hz Butterworth filter; $\sigma = 0.1$ N.
Acoustic Burst	$\mathbb{R}^1$	RMS acoustic energy envelope (100 Hz – 5 kHz).	Log-scale normalization; triggers puncture detection.
Total State $\mathcal{S}$	$\mathbb{R}^{33}$	Full Markovian observation vector fed to Actor & Critic.	Standardized via running mean and variance filter.

8.2 Continuous Action Parameterization ($\mathbb{R}^6$)

Channel	Physical Parameter	Operational Range	Robotic Function
a_0	Downward Feed Δv_z	$[-5.0, +5.0]$ mm/s	Modulates vertical penetration rate based on contact resistance.
a_1	Sawing Velocity v_{slice}	$[-30.0, +30.0]$ mm/s	Lateral sawing motion reducing normal puncture force.
a_2	Vertical Stiffness ΔK_z	$[-500, +500]$ N/m	Impedance modulation: softens blade upon initial contact to avoid crushing.
a_3	Vertical Damping ΔD_z	$[-20, +20]$ Ns/m	Damps out high-frequency contact vibrations.
a_4	Grasp Normal Force F_{hold}	$[1.0, 6.0]$ N	Holding arm normal contact force preventing tomato slip.
a_5	Holding Compliance ΔK_{hold}	$[-300, +300]$ N/m	Maintains adaptive grip on soft, deforming fruit pulp.

8.3 Multi-Objective Reward Function Formulation

The reward function balances penetration efficiency, shear mechanics, tissue damage prevention, and motion smoothness:

$$r_t = w_{\text{pen}} \cdot r_{\text{pen}} + w_{\text{slice}} \cdot r_{\text{slice}} + w_{\text{crush}} \cdot p_{\text{crush}} + w_{\text{smooth}} \cdot p_{\text{smooth}} + r_{\text{terminal}}$$

**PYTORCH TENSORIZED REWARD IMPLEMENTATION IN ISAAC LAB**

```

def compute_rewards(self) -> torch.Tensor:
    # 1. Penetration Progress: Reward downward displacement through tomato
    z_progress = torch.clamp(self.prev_knife_z - self.knife_z, min=0.0)
    r_pen = 6.0 * z_progress

    # 2. Lateral Sawing Shear: Reward sawing velocity only while blade is in contact
    v_slice = torch.abs(self.knife_vel[:, 0])
    in_contact = (self.blade_force[:, 2] > 0.4).float()
    r_slice = 3.0 * v_slice * in_contact

    # 3. Crushing Penalty: Penalize normal force exceeding 8.0 N puncture threshold
    excess_force = torch.clamp(self.blade_force[:, 2] - 8.0, min=0.0)
    p_crush = -4.0 * torch.square(excess_force)

    # 4. Action Smoothness: Penalize high-frequency chatter in impedance delta
    action_diff = self.actions - self.prev_actions
    p_smooth = -0.05 * torch.sum(torch.square(action_diff), dim=-1)

    # 5. Terminal Conditions
    r_term = torch.zeros_like(r_pen)
    # Success: Blade reaches cutting board with zero excessive crushing
    r_term = torch.where(self.reached_board & (excess_force == 0), r_term + 120.0, r_term)
    # Failure: Massive crush (>22 N) destroys the tomato
    r_term = torch.where(self.blade_force[:, 2] > 22.0, r_term - 60.0, r_term)

    return r_pen + r_slice + p_crush + p_smooth + r_term

```

8.4 Domain Randomization Strategy for Sim-to-Real Transfer

To ensure the policy trained in NVIDIA Isaac Lab transfers directly to physical Franka Emika robot arms without fine-tuning, we randomize six physical and sensory properties across 1,024 parallel environments during training:

Randomized Parameter	Nominal Value	Randomization Range	Physical Sim-to-Real Gap Targeted
Tomato Bulk Stiffness (K_s)	35 kPa	[15, 60] kPa (Uniform)	Ripeness variation between green, firm, and over-ripe tomatoes.
Skin Rupture Force (F_{rupt})	8.0 N	[5.5, 12.0] N (Gaussian)	Cuticle thickness and localized pericarp toughness variations.
Friction Coefficient ($\mu_{contact}$)	0.25	[0.12, 0.45] (Uniform)	Moisture and juice leakage lubricating the blade surface.
Blade Sharpness Index (S)	1.0 (New)	[0.6, 1.2] (Uniform)	Blade edge degradation and serration wear across repeated cuts.
Holding Gripper Compliance	800 N/m	[400, 1400] N/m	Variations in rubber finger pad elasticity and grip calibration.
Sensor Latency (Δt)	10 ms	[5, 25] ms (Stochastic)	ROS 2 network jitter and USB-FT sensor communication delays.

Module 9: Production SkRL & PyTorch Model Architecture

Below is the complete, modular PyTorch implementation of the Dual-Head Actor-Critic network configured for SkRL and Isaac Lab:

**PYTORCH DUAL-HEAD ACTOR-CRITIC MODEL (`MODELS.PY`)**

```

import torch
import torch.nn as nn
from skrl.models.torch import GaussianMixin, DeterministicMixin, Model

class SlicingActorCritic(GaussianMixin, DeterministicMixin, Model):
    def __init__(self, observation_space, action_space, device, clip_actions=False):
        Model.__init__(self, observation_space, action_space, device)
        GaussianMixin.__init__(self, clip_actions=clip_actions)
        DeterministicMixin.__init__(self)

        in_dim = observation_space.shape[0]    # 33
        out_dim = action_space.shape[0]        # 6

        # Shared representation backbone
        self.backbone = nn.Sequential(
            nn.Linear(in_dim, 256),
            nn.ELU(),
            nn.Linear(256, 256),
            nn.ELU()
        )
        # Actor head: Outputs continuous action mean
        self.actor_mean = nn.Linear(256, out_dim)
        self.log_std_parameter = nn.Parameter(torch.zeros(out_dim))

        # Critic head: Outputs scalar state-value V(s)
        self.critic_head = nn.Linear(256, 1)

    def act(self, inputs, role):
        features = self.backbone(inputs["states"])
        if role == "policy":
            mean = self.actor_mean(features)
            log_std = torch.clamp(self.log_std_parameter, min=-20.0, max=2.0)
            return mean, log_std, {}
        elif role == "value":
            value = self.critic_head(features)
            return value, {}

```

9.1 Master SkRL Configuration (`ppo\_cfg.yaml`)

Every hyperparameter in SkRL is grounded in CS285 theoretical principles:

**PRODUCTION PPO_CFG.YAML REFERENCE**

```

algorithm:
  class: PPO
  clip_range: 0.2           # Lecture 10: Enforces trust region, preventing policy collapse
  gae_lambda: 0.95         # Lecture 6: Optimal exponential blend of TD(0) bias and MC variance
  discount_factor: 0.99    # Lecture 4:  $\gamma = 0.99$  ensures credit assigned across 300+ cutting steps
  learning_rate: 3e-4       # Standard Adam step size
  learning_rate_scheduler: KLLAdaptiveLR # Adaptive step size scaling based on KL divergence
  entropy_loss_scale: 0.01  # Lecture 8: Encourages exploration of dynamic impedance angles
  value_loss_scale: 1.0     # Weight of Critic MSE regression loss
  epochs: 5                # Lecture 10: Re-uses parallel rollout batches safely
  mini_batches: 4          # Subdivides 1,024 parallel environments into mini-batches

```


Module 10: Thesis Defense Master Cheatsheet (Top 10 Q&A)

When defending your methodology before Prof. Shan An, thesis committees, or IEEE reviewers, be prepared for these 10 core theoretical questions:

#	Reviewer / Committee Question	Your Bulletproof Academic Answer
1	Why use Reinforcement Learning instead of standard Impedance Control?	Classical impedance control requires hand-tuning constant stiffness (K) and damping (D). Soft fruits exhibit non-linear viscoplastic fracture transitions where contact stiffness drops by over 80% within milliseconds upon skin puncture. RL dynamically modulates compliance parameters in response to real-time high-frequency tactile and acoustic signals.
2	Why is your task formulated as a Markov Decision Process (MDP)?	Our 33-dimensional state vector contains both kinematics (joint and blade positions/velocities) and contact forces (F/T), meaning the transition to s_{t+1} depends solely on the current state s_t and action a_t , strictly satisfying the Markov conditional independence property.
3	Why did you choose PPO over Model-Based RL?	Deformable viscoplastic fracture mechanics are notoriously difficult to model analytically. Model-based RL compounds model bias across extended horizons, causing sim-to-real transfer failure. PPO is model-free, highly robust, and directly optimizes the true multi-objective reward.
4	How do you address PPO's low sample efficiency?	Through GPU-vectorized parallel physics simulation in NVIDIA Isaac Lab. By simulating 1,024 environments simultaneously on an RTX GPU, we collect over 10 million simulation transitions in under 45 minutes, rendering on-policy sample efficiency irrelevant.
5	What is the exact role of the Critic network?	The Critic approximates the State-Value function $V(s)$ using Temporal Difference learning. It is used to compute the Advantage function $\hat{A}(s, a) = Q(s, a) - V(s)$, which dramatically reduces policy gradient variance compared to raw Monte Carlo rollouts.
6	Why is Generalized Advantage Estimation (GAE) with $\lambda = 0.95$ used?	Pure TD learning ($\lambda = 0$) has minimal variance but introduces bias if the Critic is inaccurate. Pure rollouts ($\lambda = 1$) have zero bias but massive variance. GAE with $\lambda = 0.95$ achieves an optimal exponential blend of multi-step returns.
7	Why did you configure continuous action spaces instead of discrete?	Robot joint motor torques, feed velocities, and compliance deltas ($\Delta K, \Delta D$) are continuous physical quantities. Discretizing them creates coarse, jerky motions that damage soft object skins.
8	What is the purpose of PPO's clipping mechanism?	It bounds the probability ratio $r_t(\theta)$ between $[1 - \epsilon, 1 + \epsilon]$ (typically $\epsilon = 0.2$). This prevents policy collapse by discarding gradient steps that deviate too far from the policy that generated the trajectory data.
9	Why is lateral sawing velocity rewarded?	Under fracture mechanics shear-stress coupling, lateral slicing motion significantly reduces the required normal downward force (F_z) to achieve fracture toughness (K_{Ic}), preventing fruit crushing.
10	How does Soft Actor-Critic (SAC) prevent local minima?	By optimizing Maximum Entropy RL (Reward + $\alpha \cdot$ Entropy). The entropy term encourages broad, stochastic exploration, preventing the robot from adopting timid policies that freeze above the tomato skin to avoid crushing penalties.

DEX-ROB Lab (Tianjin University) Master's Research Program • Prof. Shan An • Created with Antigravity AI